

# MR 典型编程场景 2

## 目录

1、自定义 OutputFormat--数据分类输出 .....	1
1.1、需求.....	1
1.2、分析.....	1
1.3、实现.....	1
2、自定义 InputFormat--小文件合并 .....	5
2.1、需求.....	5
2.2、分析.....	5
2.3、实现.....	5

## 1、自定义 OutputFormat--数据分类输出

### 1.1、需求

现有一些原始日志需要做增强解析处理，流程：

- 1、从原始日志文件中读取数据
- 2、根据业务获取业务数据库的数据
- 3、根据某个连接条件获取相应的连接结果

典型业务场景如：爬虫 URL 管理，移动号码管理

### 1.2、分析

程序的关键点是要在一个 MapReduce 程序中根据数据的不同输出两类结果到不同目录，这类灵活的输出需求可以通过自定义 OutputFormat 来实现

### 1.3、实现

实现要点：

- 1、在 MapReduce 中访问外部资源
- 2、自定义 OutputFormat，改写其中的 RecordWriter，改写具体输出数据的方法 write()

以 Score.txt 的 32 条学生考试记录为例，现要求把参考次数 $\geq 7$  的输出到一个文件/output/out1，然后剩下的不合格的参考输出到另外一个文件/output/out2

实现:

第一步, 实现自己的 OutputFormat

```
package com.ghgj.mr.format.output;

import java.io.IOException;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FSDataOutputStream;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IOUtils;
import org.apache.hadoop.io.NullWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.RecordWriter;
import org.apache.hadoop.mapreduce.TaskAttemptContext;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class LongOutputFormat extends FileOutputFormat<Text, NullWritable>{

    @Override
    public RecordWriter<Text, NullWritable> getRecordWriter(
        TaskAttemptContext job) throws IOException, InterruptedException {

        Configuration configuration = job.getConfiguration();
        FileSystem fs = FileSystem.get(configuration);

        Path p1 = new Path("d:/outputformat/out1");
        Path p2 = new Path("d:/outputformat/out2");

        FSDataOutputStream out1 = fs.create(p1);
        FSDataOutputStream out2 = fs.create(p2);

        return new MyRecordWriter(out1, out2);
    }

    static class MyRecordWriter extends RecordWriter<Text, NullWritable>{

        FSDataOutputStream fsdout = null;
        FSDataOutputStream fsdout1 = null;

        public MyRecordWriter(FSDataOutputStream fsdout,
            FSDataOutputStream fsdout1) {
            super();
            this.fsdout = fsdout;
        }
    }
}
```

```
        this.fsdout1 = fsdout1;
    }

    @Override
    public void write(Text key, NullWritable value) throws IOException,
        InterruptedException {
        String[] str = key.toString().split("::");
        if(strs[0].equals("1")){
            fsdout.write((strs[1]+"\\n").getBytes());
        }else{
            fsdout1.write((strs[1]+"\\n").getBytes());
        }
    }

    @Override
    public void close(TaskAttemptContext context) throws IOException,
        InterruptedException {
        IOUtils.closeStream(fsdout);
        IOUtils.closeStream(fsdout1);
    }
}
```

## 第二步，实现 MapReduce 程序

```
package com.ghgj.mr.format.output;

import java.io.IOException;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.NullWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class MultipleOutputMR {

    public static void main(String[] args) throws Exception {

        Configuration conf = new Configuration();
```

```
Job job = Job.getInstance(conf);

job.setJarByClass(MultipleOutputMR.class);

job.setMapperClass(MultipleOutputMRMapper.class);
job.setMapOutputKeyClass(Text.class);
job.setMapOutputValueClass(NullWritable.class);

job.setOutputFormatClass(LongOutputFormat.class);

FileInputFormat.setInputPaths(job, "d:/score/input");

Path outputPath = new Path("d:/score/output_success");
FileSystem fs = FileSystem.get(conf);
if (fs.exists(outputPath)) {
    fs.delete(outputPath, true);
}
FileOutputFormat.setOutputPath(job, outputPath);

boolean waitForCompletion = job.waitForCompletion(true);
System.exit(waitForCompletion ? 0 : 1);
}

static class MultipleOutputMRMapper extends
    Mapper<LongWritable, Text, Text, NullWritable> {

    @Override
    protected void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {

        // value
        // 参考次数大于 7 次算合格
        String[] splits = value.toString().split("\t");
        if (splits.length > 9) {
            context.write(new Text("1::" + value.toString()), NullWritable.get());
        } else {
            context.write(new Text("2::" + value.toString()), NullWritable.get());
        }
    }
}
}
```

## 2、自定义 InputFormat--小文件合并

### 2.1、需求

无论 HDFS 还是 MapReduce，对于小文件都有损效率，实践中，又难免面临处理大量小文件的场景，此时，就需要有相应解决方案

### 2.2、分析

小文件的优化无非以下几种方式：

- 1、在数据采集的时候，就将小文件或大批数据合成大文件再上传 HDFS
- 2、在业务处理之前，在 HDFS 上使用 MapReduce 程序对小文件进行合并
- 3、在 MapReduce 处理时，可采用 CombineFileInputFormat 提高效率

### 2.3、实现

在此，我们采用第二种方式使用 MapReduce 程序来对小文件进行合并。注意：并不是说编写一个 MR 程序来实现对这小文件的计算，只是做合并

核心实现思路：

- 1、编写自定义的 InputFormat
- 2、改写 RecordReader，实现一次 maptask 读取一个小文件的完整内容封装了一个 KV 对
- 3、在 Driver 类中一定要设置使用自定义的 InputFormat：  
`job.setInputFormatClass(WholeFileInputFormat.class)`

看具体实现：

第一步，编写自定义的 InputFormat

```
package com.ghgj.mr.format.input;

import java.io.IOException;

import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.NullWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.InputSplit;
import org.apache.hadoop.mapreduce.JobContext;
import org.apache.hadoop.mapreduce.RecordReader;
import org.apache.hadoop.mapreduce.TaskAttemptContext;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;

public class WholeFileInputFormat extends FileInputFormat<NullWritable, Text> {
```

```
// 设置每个小文件不可分片,保证一个小文件生成一个 key-value 键值对
@Override
protected boolean isSplittable(JobContext context, Path file) {
    return false;
}

@Override
public RecordReader<NullWritable, Text> createRecordReader(InputSplit split,
TaskAttemptContext context) throws IOException, InterruptedException {
    WholeFileRecordReader reader = new WholeFileRecordReader();
    reader.initialize(split, context);
    return reader;
}
}
```

## 第二步，编写自定义的 RecordReader

```
package com.ghgj.mr.format.input;

import java.io.IOException;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FSDataInputStream;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IOUtils;
import org.apache.hadoop.io.NullWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.InputSplit;
import org.apache.hadoop.mapreduce.RecordReader;
import org.apache.hadoop.mapreduce.TaskAttemptContext;
import org.apache.hadoop.mapreduce.lib.input.FileSplit;

class WholeFileRecordReader extends RecordReader<NullWritable, Text> {
    private FileSplit fileSplit;
    private Configuration conf;
    private Text value = new Text();
    private boolean processed = false;

    @Override
    public void initialize(InputSplit split, TaskAttemptContext context)
        throws IOException, InterruptedException {
        this.fileSplit = (FileSplit) split;
        this.conf = context.getConfiguration();
    }
}
```

```
@Override
public boolean nextKeyValue() throws IOException, InterruptedException {
    if (!processed) {
        byte[] contents = new byte[(int) fileSplit.getLength()];
        Path file = fileSplit.getPath();
        FileSystem fs = file.getFileSystem(conf);
        FSDataInputStream in = null;
        try {
            in = fs.open(file);
            // 把输入流上的数据全部读取到 contents 字节数组里
            IOUtils.readFully(in, contents, 0, contents.length);
            // 把读取到的数据设置到 value 里
            value.set(contents, 0, contents.length);
        } finally {
            IOUtils.closeStream(in);
        }
        processed = true;
        return true;
    }
    return false;
}

@Override
public NullWritable getCurrentKey() throws IOException, InterruptedException {
    return NullWritable.get();
}

@Override
public Text getCurrentValue() throws IOException, InterruptedException {
    return value;
}

@Override
public float getProgress() throws IOException {
    return processed ? 1.0f : 0.0f;
}

@Override
public void close() throws IOException {
    // do nothing
}
}
```

### 第三步，编写 MapReduce 程序

```
package com.ghgj.mr.format.input;

import java.io.IOException;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.NullWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.InputSplit;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.FileSplit;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.util.Tool;
import org.apache.hadoop.util.ToolRunner;

public class SmallFilesConvertToBigMR extends Configured implements Tool {

    public static void main(String[] args) throws Exception {
        int exitCode = ToolRunner.run(new SmallFilesConvertToBigMR(), args);
        System.exit(exitCode);
    }

    static class SmallFilesConvertToBigMRMapper extends Mapper<NullWritable, Text, Text, Text> {
        private Text filenameKey;

        @Override
        protected void setup(Context context) throws IOException, InterruptedException {
            InputSplit split = context.getInputSplit();
            Path path = ((FileSplit) split).getPath();
            filenameKey = new Text(path.toString());
        }

        @Override
        protected void map(NullWritable key, Text value, Context context)
            throws IOException, InterruptedException {
            context.write(filenameKey, value);
        }
    }
}
```



```
static class SmallFilesConvertToBigMRReducer extends Reducer<Text, Text, NullWritable, Text> {
    @Override
    protected void reduce(Text filename, Iterable<Text> bytes,
        Context context) throws IOException, InterruptedException {
        context.write(NullWritable.get(), bytes.iterator().next());
    }
}

@Override
public int run(String[] args) throws Exception {
    Configuration conf = new Configuration();

    conf.set("fs.defaultFS", "hdfs://hadoop02:9000");
    System.setProperty("HADOOP_USER_NAME", "hadoop");

    Job job = Job.getInstance(conf, "combine small files to bigfile");
    job.setJarByClass(SmallFilesConvertToBigMR.class);

    job.setMapOutputKeyClass(Text.class);
    job.setMapOutputValueClass(Text.class);
    job.setMapperClass(SmallFilesConvertToBigMRMapper.class);

    job.setReducerClass(SmallFilesConvertToBigMRReducer.class);
    job.setOutputKeyClass(NullWritable.class);
    job.setOutputValueClass(Text.class);

    job.setInputFormatClass(WholeFileInputFormat.class);
    // job.setOutputFormatClass(SequenceFileOutputFormat.class);

    Path input = new Path("/smallfiles");
    Path output = new Path("/bigfile");
    FileInputFormat.setInputPaths(job, input);
    FileSystem fs = FileSystem.get(conf);
    if (fs.exists(output)) {
        fs.delete(output, true);
    }
    FileOutputFormat.setOutputPath(job, output);

    int status = job.waitForCompletion(true) ? 0 : 1;
    return status;
}
```